BUSINESS PROPOSAL

# Ronda.

*An AI shift coach for Malaysian gig riders.*

Domain 2 · AI for Economic Empowerment & Decision Intelligence

UMHackathon 2026 · Final Round · 3 May 2026


## Team Ronda

Srinivaas · Thenesh · Teeva · Davud · Ambrish


*Powered by ILMU-GLM-5.1*

in partnership with Z.AI and YTL AI Labs

01

# Executive Summary

> *Every morning, 1.2 million Malaysian gig workers make the same RM-significant decision blind: should I work today, where, and when? No tool answers it. We built one.*

Ronda is an AI shift coach for Malaysian gig riders. It synthesises five fragmented inputs — weather, platform incentives, local events, fuel-price quota status, and the rider's own earnings history — into a single structured shift recommendation: a verdict (work / rest / partial), shift windows with zones, projected net earnings after petrol, and a Malaysian-English narrative that reads like a friend who actually knows the numbers. The reasoning engine is ILMU-GLM-5.1, accessed through an Anthropic-compatible API surface. Removing it degrades Ronda to a generic rules baseline that cannot read the calendar, weight competing incentives, or calibrate its own confidence — a gap the product makes visible at /compare.

### The opportunity, in three numbers

| 1.2M | RM 2,185 | 98% |
|---|---|---|
| Malaysian gig workers (Gig Workers Act 2025) | Avg rider monthly income — 50% below national | Riders aged 15–30 (DOSM) |

The market is large, young, economically strained, and structurally underserved. As of April 2026, Malaysia's Gig Workers Act 2025 (Act 872) is active law, but PERKESO confirms only 2% of the 1.2 million eligible workers have been onboarded so far. Fuel subsidy rationalisation has put petrol cost variance into every rider's daily decision. The platforms hide data because asymmetry is their margin. Nobody has built rider-side decision intelligence. Now is the open window.

### The business in one paragraph

Ronda is a B2C-acquired, B2B-monetised decision intelligence platform. Riders use it free, generating proprietary behavioural data in exchange. We monetise three layers: a free tier that drives organic acquisition through WhatsApp screenshots, a Ronda+ subscription at RM 9 per month for power users, and Ronda Intelligence — anonymised aggregate insights sold to insurers, fuel brands, PERKESO compliance teams, and policymakers who currently have no neutral source on the gig economy. The riders are not the customer. The intelligence they generate is the product.

## Why this is fundable

Three moats compound from day one. A proprietary behavioural dataset that thickens with every signup. Trust capital — riders share Ronda screenshots in WhatsApp groups, they don't share platform incentive bulletins. And alignment with Malaysian sovereign AI infrastructure: Ronda is among the first products built on the ILMU AI Accelerator Programme, jointly run by YTL AI Labs and MDEC. Government and institutional partners explicitly prefer sovereign infrastructure for data-sensitivity reasons. This is a feature for procurement, not a constraint.

## Traction

Ronda has shipped a working product. The web application is live, integrated end-to-end with ILMU-GLM-5.1, validated by a Pydantic schema with a 100% first-try pass rate, and protected by a three-tier fallback chain that maintains availability under upstream degradation. Team Ronda won the UMHackathon 2026 preliminary round and is presenting in the national finals on 3 May 2026.

## The ask

This document is a thesis, not a fundraise. We are presenting a category-creating product, a credible 18-month operating plan, and a path to RM 16.7 million in revenue by Year 3 in Malaysia alone — before any regional expansion.

02

# The Product

Ronda is a single-screen web application. A rider opens it, optionally enters today's context (location, hours available, fuel-quota status, weather), and receives a structured shift recommendation in approximately ten seconds. The output is designed to be glanceable on a phone before logging into the platform apps.

### What a rider sees

- **A verdict.** Work, rest, or partial. One word, foregrounded.
- **Shift windows with zones.** When to be online and where to position. Calibrated to the rider's home zone, not generic city-wide advice.
- **Projected net earnings after petrol.** Not gross. Net. Because gross is the lie the platforms tell.
- **Ranked key factors.** Weather, surge expectation, event-driven demand, incentive structure, fuel-quota burn rate. Each weighted, each explained.
- **Explicit caveats.** When confidence is low, Ronda says so in plain Malaysian English. No black box, no false certainty.

### What makes the engineering credible

- **Three-tier reliability.** Streaming GLM is the preferred path. On Cloudflare timeout or stream drop, the system falls back to non-streaming. On full upstream failure, a rule-based baseline grounded in the rider's 7-day average activates, tagged source: "fallback_rules" so the UI communicates degraded mode honestly.
- **Context compression.** A custom summariser reduces 3,200-character ContextPackets to 350-character paragraphs before they reach GLM, cutting latency from a 504 timeout to ten seconds while preserving decision-relevant signal.
- **Schema-enforced output.** Every GLM response is parsed against a Pydantic schema. Drift triggers a corrective retry that feeds the validation error back to GLM. First-try pass rate measured at 100% on production inputs.
- **Side-by-side compare view.** /compare displays GLM's reasoning next to a rules-only baseline using identical input. This is both a debugging tool and our demo's wow moment: judges and users see, in real time, why GLM is non-substitutable.

### What GLM does that rules cannot

Removing the reasoning engine degrades Ronda in four specific, demonstrable ways.

- **Event awareness.** Rules cannot read "Coldplay at Bukit Jalil tonight" from a calendar feed and infer surge implications across e-hailing demand, food delivery on the perimeter, and stadium-traffic congestion windows. GLM does this in one pass.

- **Incentive weighting.** When Grab offers RM30 for 8 trips and Foodpanda offers RM50 for 12 trips on the same day, the answer depends on weather, the rider's batch level, fuel quota remaining, and shift window. Rules collapse into a heuristic. GLM reasons.
- **Confidence calibration.** GLM tells the rider how sure it is and why, in language they trust. Rules are either overconfident or silent.
- **Cultural context.** GLM understands that Friday afternoon during Ramadan is not a normal Friday afternoon. That Thaipusam at Batu Caves shifts the entire northern Klang Valley demand pattern. That a rainy Monday in Bangsar is different from a rainy Monday in Cheras. ILMU's Malay-MMLU-leading Bahasa Melayu and Manglish fluency makes this work natively.

None of these are bolt-on features. They are emergent capabilities of the reasoning engine. The /compare view exists in production because we want this delta to be visible, not asserted.

03

# Market Research & Dynamics

## 3.1 The Problem

The Malaysian gig economy has scaled faster than the support infrastructure around it. The Department of Statistics Malaysia tracked 4 million workers in the broader gig economy as of 2023, with 3.25 million own-account workers as of October 2025. The Gig Workers Act 2025, which came into force on 31 March 2026, identifies 1.2 million workers in the platform-mediated subset — primarily food delivery, p-hailing, and e-hailing — concentrated heavily in Klang Valley, Penang, and Johor Bahru.

Of these workers, 98% are aged 15 to 30. 73% rely on gig work as their main source of income. The average Foodpanda rider in Malaysia earns RM 2,185 per month — 50% below the national median formal-sector wage of RM 3,000. Foodpanda's batch system pays RM 4.50 to RM 7.00 per order depending on rank, producing weekly earnings between RM 396 and RM 600 for ten-hour days. In January 2024, riders staged a 24-hour "Food Delivery Blackout" strike protesting unsustainable wages. The 2024 "Grab Blackout" followed. Public discontent is documented and structural, not anecdotal.

> *It's frustrating that I earn very little from my rider fees despite travelling to many places to deliver orders. I work for more than 10 hours per day, and it gets even more stressful when I get scolded by customers for delays caused by vendors.*
>
> — Mohd Khairul Affendy, Foodpanda rider, quoted in Scoop, January 2024

Every one of these workers makes the same decision every day: should I go out, where, and when. The decision determines whether they earn RM 200 or RM 40, whether they pay rent on time, whether the day was worth the petrol. And every one of them makes it blind.

The data they need exists, but it lives in five separate places. Weather is in a weather app. Incentive bulletins are buried in the platform driver app. Concerts and stadium events are in news feeds and Telegram. Fuel prices, BUDI95 quota status, and unsubsidised market rates update weekly. Their own earnings history is in a spreadsheet, on paper, or in their head. No tool stitches these together — not because it is technically hard, but because the platforms benefit from the asymmetry, and no neutral party has had the incentive to do it.

> *The platforms know what every rider should do. They just have no reason to tell them.*

## 3.2 Why This Is The Right Moment

Three forces are pushing this market open right now, simultaneously.

### Regulatory tailwind

The Gig Workers Act 2025 (Act 872) became active law on 31 March 2026. It mandates SOCSO contributions, written service agreements, dispute tribunals, and explicit prohibition of unilateral rate changes by platforms. As of late April 2026, PERKESO publicly stated that only 2% of the 1.2 million eligible workers had been onboarded. The compliance gap is enormous, the infrastructure for tracking it is absent, and no neutral data source exists. This is a classic regulatory window: legislation exists but operational tooling does not.

### Fuel cost volatility

The BUDI95 programme caps subsidised RON95 at RM 1.99 per litre, with an 800-litre monthly quota for e-hailing and gig workers. Beyond quota, riders pay floating market rates — currently around RM 2.60 per litre and projected by AMRO to reach RM 2.90 if subsidies are floated to market. Fuel cost is no longer a fixed background expense. It is now a daily decision input — exactly what Ronda models in its net-earnings projection.

### Sovereign AI is suddenly viable

ILMU was launched by YTL AI Labs in August 2025 as Malaysia's first homegrown multimodal LLM, in collaboration with Universiti Malaya. It outperforms GPT-4o, Llama 3.1 and DeepSeek V3 on the Malay-MMLU benchmark. The ILMU AI Accelerator Programme, run jointly with MDEC, offers Malaysia Digital status, USD 1 million in API credits, and talent-pass facilitation to startups building on the platform. Until 2025, building a Malaysian-context reasoning product meant fine-tuning a Western model at high cost or accepting confident hallucination of local geography. That is no longer the trade-off. Ronda is a first mover on this infrastructure.

## 3.3 Market Sizing

All figures grounded in DOSM, PERKESO, and Ministry of Human Resources data. Methodology disclosed below the table.

| Tier | Definition | Estimated Value (Year 3) |
|---|---|---|
| TAM | Malaysian gig worker decision-intelligence market — all 1.2 million platform-mediated workers identified under the Gig Workers Act 2025, across food delivery, p-hailing, e-hailing. | RM 1.44 billion |
| SAM | Active full-time and serious part-time riders in Klang Valley, Penang, and Johor Bahru — the three urban tiers where Ronda's recommendations have material zone-level resolution. ~ 410,000 workers. | RM 492 million |

| Tier | Definition | Estimated Value (Year 3) |
|---|---|---|
| SOM | Realistic 18-month capture: 12% of Klang Valley active riders on free tier, 9% conversion to Ronda+, plus four B2B Intelligence contracts. | RM 72 million ARR |

***Methodology.*** *TAM derived from 1.2M Gig Workers Act-eligible workers × blended average annual rider value (RM 9/month rider tier ARPU × estimated paid penetration + B2B intelligence value per worker). SAM narrows to ~ 34% concentration in three urban tiers per DOSM regional employment data. SOM reflects an 18-month conservative capture in Klang Valley first, with B2B contract pricing benchmarked against comparable data products (Strava Metro, Waze for Cities). Sources: DOSM Labour Force Statistics (Oct 2025), PERKESO press releases (April 2026), Ministry of Human Resources Gig Workers Act briefing (Sept 2025), Foodpanda rider income data via Scoop and Indeed Malaysia.*

04

# Value & Moat

## 4.1 The Three Antagonists

Ronda's positioning becomes legible once the antagonists are named. There are three.

### The platforms

Grab, Foodpanda, ShopeeFood, and the e-hailing operators benefit from rider information asymmetry. They know surge timing, they know which zones underperform, they know which riders are about to churn. They will not share this with the rider, because the rider's information disadvantage is the platform's negotiating leverage. The Gig Workers Act 2025 explicitly prohibits unilateral rate changes — the legislation exists because the conduct exists. Ronda is structurally on the rider's side, in a way the platforms cannot be.

### The chaos

Five apps, five tabs, five sources of truth. Weather, incentives, news, fuel, and the rider's own scattered records. No synthesis. No narrative. No verdict. The cost of the chaos is not abstract — it is the difference between a RM 200 day and a RM 60 day, and we have rider testimony confirming that gap is routine.

### The intuition

In the absence of tools, riders rely on gut feel and informal WhatsApp sharing. This is free, fast, and culturally entrenched — and it is wrong often enough to cost the average rider an estimated RM 50 to 150 per month. Intuition is the substitute we beat by being demonstrably better. The /compare view exists in part to make this delta visible, not asserted.

## 4.2 What Makes This Defensible

Three moats. They compound.

### The data flywheel

Every rider who uses Ronda generates a labelled behavioural data point: this rider, in this zone, on this day, with these conditions, made this decision and earned this much. After twelve months, Ronda owns the only longitudinal dataset of Malaysian gig-worker decisions outside the platforms themselves. After twenty-four months, it is the most comprehensive. This dataset is the B2B product. It is also the model improvement loop — every interaction makes the next recommendation sharper.

### Trust capital

Riders share Ronda screenshots in WhatsApp groups today. They do not share Grab's incentive bulletin. The reason is structural: Ronda is on their side, the platform is not. This trust is asymmetric and not

replicable by the platforms, who would be giving up their information advantage to compete. It is also not replicable by a foreign entrant — Telegram channels took two years to earn rider trust, and they did so by being run by riders. Ronda inherits that register through its launch strategy (community-led, not paid-acquisition-led).

### Sovereign AI alignment

Ronda is built on ILMU-GLM-5.1, a Malaysian-context reasoning engine that no Silicon Valley competitor can match without a years-long localisation effort. Government and institutional procurement explicitly prefer sovereign infrastructure for data-sensitivity reasons — the ILMU AI Accelerator Programme run by YTL AI Labs and MDEC exists precisely to channel demand toward Malaysian-built products. For Ronda's B2B Intelligence buyers (insurers, government agencies, PERKESO compliance teams), sovereign AI is not a marketing line. It is a procurement requirement.

> *Our moat is not the AI. The AI is accessible to anyone. Our moat is the proprietary dataset of Malaysian gig-worker behaviour that compounds with every user. Nobody else has it. Grab does, but they will never sell it. We become the only neutral source.*

05

# Competitive Landscape

There are no direct competitors today. No product synthesises rider-side data into a daily decision in Malaysia. The relevant analysis is therefore about substitution risk and adjacent forces.

## 5.1 Porter's Five Forces

| Force | Intensity | Assessment |
|---|---|---|
| Competitive Rivalry | Low | No direct competitor builds rider-side decision intelligence in Malaysia. Adjacent players (Telegram channels, Facebook groups, freelance YouTube guides) are unstructured and not personalised. Ronda is creating, not entering, the category. |
| Threat of New Entrants | Medium | Technical barrier is moderate — a competent team could build v1 in three months. Data barrier compounds: by month twelve, Ronda's behavioural dataset is hard to replicate without comparable rider trust. Trust barrier is highest: Telegram channels took two years to earn theirs. |
| Bargaining Power of Riders | Low | Switching cost from intuition to Ronda is zero. From Ronda to a hypothetical competitor would require trust transfer and historical data loss. Free tier removes the price-sensitivity barrier. |
| Bargaining Power of B2B Buyers | Medium | Insurers and fuel brands have moderate power individually but no alternative source of equivalent data. Government agencies (MDEC, PERKESO, DOSM) need this data for policy enforcement and have budget to pay for it. |
| Bargaining Power of Suppliers | Medium | Primary supplier is the ILMU-GLM-5.1 partnership. Anthropic-compatible API surface allows multi-provider fallback technically, but loss of Malaysian context would degrade the product. The three-tier fallback chain already proves Ronda functions in degraded mode. |
| Threat of Substitutes | Medium | Biggest substitute is the rider's own intuition plus informal WhatsApp sharing. Free, ambient, culturally entrenched. Ronda must beat informal channels on three dimensions: speed, accuracy, trust. /compare exists to make this delta visible. |

## 5.2 Comparison

| Competitor | Strength | Weakness | Ronda's edge |
|---|---|---|---|
| Platform driver apps (Grab, Foodpanda) | Native distribution. Already in every rider's pocket. | Structurally conflicted — cannot tell a rider to switch off without harming their own marketplace. | Ronda is rider-side. We win on alignment, not features. |
| Rider Telegram channels | High trust, fast, free, culturally native. | Unstructured, unverifiable, no personalisation, no historical data. | Structured personalisation while preserving the cultural register that built that trust. |
| Generic AI chatbots (ChatGPT, Gemini, Copilot) | Powerful general reasoning. | No Malaysian context, no live data integration, hallucinated geography. | Built on Malay-MMLU-leading ILMU with explicit fallback when uncertain. |
| Manual intuition (the default) | Free, instant, no app needed. | Costs the average rider an estimated RM 50–150 per month in suboptimal decisions. | Makes the cost of intuition visible, then beats it. |

06

# Business Model

## 6.1 Customer Segments

- **Active Klang Valley Rider — Primary.** Full-time gig rider, RM 2,000–RM 4,500/month, one to three years of platform experience, multi-platform, smartphone-fluent, WhatsApp-active. The Aiman persona. Converts to Ronda+ at the highest rate.
- **Part-time and Weekend Rider — Secondary.** Supplemental income, lower frequency, harder to convert to paid but valuable as a data contributor and viral-acquisition node. The Mei Ling persona.
- **Multi-Platform Power User — High-value.** Runs three or more platforms simultaneously, optimises aggressively, currently builds personal spreadsheets. Highest willingness-to-pay. The Ravi persona.
- **Institutional Buyer — B2B.** Gig-worker insurance underwriters (Etiqa, Allianz, Zurich), fuel brands (Petronas, Shell), telco prepaid teams, and government agencies (MDEC, MOHR, PERKESO, DOSM). These customers buy Ronda Intelligence.

## 6.2 Revenue Streams

Ronda runs a three-layer model. The riders are not the customer; the data they generate is the product.

| Layer | Pricing | Included | Strategic role |
|---|---|---|---|
| Free | RM 0 | Daily morning brief, one recommendation per day, basic narrative. | Acquisition engine. Drives organic WhatsApp-screenshot virality and feeds the data flywheel. |
| Ronda+ | RM 9 / month | Unlimited recommendations, multi-day planning, earnings tracking, weather radar overlay, BUDI95 quota tracker, priority GLM tier. | Direct revenue. Priced below one hour of rider earnings to keep the value math obvious. |
| Ronda Intelligence | RM 60K–500K / contract / year | Anonymised aggregate behavioural data, custom dashboards, API access, market reports, PERKESO | The margin engine. This is what the company is actually worth. |

| Layer | Pricing | Included | Strategic role |
|---|---|---|---|
| | | compliance views. | |

### The Strava parallel

This is the Strava model and the Waze model, ported to Malaysian gig work. Strava gave cyclists free tracking, then sold aggregated bike-flow data to cities planning infrastructure (Strava Metro). Waze gave drivers free navigation, then sold traffic intelligence to municipalities and brands (Waze for Cities). Both companies became defensible monopolies on a dataset they got users to volunteer. Ronda follows the same playbook on a market with sharper economic urgency, weaker existing tooling, and an active regulatory regime that needs the data we generate.

## 6.3 Cost Structure

| Cost Driver | Type |
|---|---|
| ILMU-GLM-5.1 inference (per recommendation) | Variable — scales with usage; mitigated by context compression and ILMU AI Accelerator credits |
| Cloud infrastructure (FastAPI hosting, database, CDN) | Fixed → Variable as usage scales |
| Engineering and product team | Fixed (largest cost in Year 1) |
| Rider acquisition (referral incentives, content, community) | Variable — highest in months 1–6, decreasing as virality compounds |
| B2B sales and partnership development | Fixed — one senior hire from month 9 onwards |
| Data infrastructure (warehouse, anonymisation, BI) | Fixed → Variable as B2B contracts scale |

## 6.4 Key Partners

| Strategic | Operational |
|---|---|
| YTL AI Labs (ILMU model partner via AI Accelerator) | Cloud infrastructure (AWS Asia Pacific or YTL AI Cloud) |

| Strategic | Operational |
|---|---|
| MDEC (Malaysia Digital status, accelerator alignment) | Payment rails (Stripe, Billplz, FPX) |
| Insurance underwriters (first B2B Intelligence customers) | Weather data (Malaysian Meteorological Department feeds) |
| Rider community leaders & Telegram admins (distribution) | Events data (Klook, BookMyShow, news aggregators) |
| PERKESO (compliance data partnership) | BUDI95 quota status (KPDN data feeds) |

## 6.5 Risk Assessment

| Risk | Impact | Likelihood | Mitigation |
|---|---|---|---|
| Platforms (Grab/Foodpanda) build a competing first-party tool | High | Medium | Structural conflict: any platform tool will favour the platform's marketplace. Ronda's neutrality is the moat. Position publicly as rider-side from day one. |
| GLM API outage or breaking version change | High | Low | Three-tier fallback chain (streaming → non-streaming → rule-based) already shipped. Anthropic-compatible surface allows multi-provider fallback if required. |
| Slow rider conversion to paid tier | Medium | Medium | B2B Intelligence revenue is the primary margin engine; rider paid is supplementary. Free tier still produces data with no conversion required. |
| PDPA / regulatory restriction on aggregation | High | Low to Medium | Anonymisation pipeline designed for PDPA compliance from day one. Government agencies (MDEC, PERKESO, DOSM) modelled as customers, not adversaries — early policy engagement de-risks this entirely. |
| GLM hallucination causes rider to lose money | Medium | Low | Schema validation, corrective retry on drift, explicit confidence calibration in output, published methodology page. Users see caveats, not a black box. |

07

# Implementation & Financials

## 7.1 Go-to-Market

Ronda's go-to-market is sequenced around a single insight: in the Malaysian gig economy, distribution is built in WhatsApp groups and Telegram channels, not on Instagram and not on billboards. The product launches into communities that already exist.

Phase 1 is a Klang Valley-only launch. We seed Ronda into 30 of the largest rider-only WhatsApp and Telegram groups through community leaders we have already mapped, paired with a referral incentive (RM 10 e-wallet credit per rider activation, capped). Screenshots of the morning brief are designed to be screenshot-shareable — pull-quote typography, no UI chrome — converting daily-use into organic distribution.

Phase 2 expands to Penang and Johor Bahru once unit economics in Klang Valley are proven (target: CAC under RM 12, D30 retention above 35%). The B2B sales motion begins in parallel with a single senior partnerships hire targeting two pilot Intelligence contracts in Year 1 — most plausibly with an insurance underwriter and PERKESO.

Phase 3 opens conversations with regional partners (Indonesia, Vietnam, Philippines) using the proven Malaysian playbook. Geographic expansion is explicitly out of scope for the 18-month plan below.

## 7.2 Implementation Plan

| Phase | Action | Timeline | Target |
|---|---|---|---|
| Soft Launch | Klang Valley only. Seed via 30 WhatsApp/Telegram groups. Daily content cadence on rider-economic stories. | Months 1–3 | 8,000 free MAU, 600 Ronda+ subscribers |
| Growth | Referral incentive engine, paid acquisition test, first B2B Intelligence pilot signed. | Months 4–9 | 35,000 free MAU, 3,200 Ronda+, 1 B2B pilot live |
| Multi-City | Penang and Johor Bahru launches. Second B2B contract. Senior partnerships hire. | Months 10–14 | 85,000 free MAU, 7,800 Ronda+, 2 B2B contracts |
| Scale | Optimise unit economics. Open SEA | Months 15–18 | 140,000 free MAU, 13,500 |

| Phase | Action | Timeline | Target |
|---|---|---|---|
| | conversations. Begin Series A preparation. | | Ronda+, 4 B2B contracts |

## 7.3 Revenue Projections

All figures in MYR. Projections reflect Klang Valley plus two-city Malaysian rollout only. No regional revenue is modelled.

| Metric | Year 1 | Year 2 | Year 3 |
|---|---|---|---|
| Active free MAU | 35,000 | 120,000 | 240,000 |
| Ronda+ subscribers | 3,200 | 12,500 | 28,000 |
| B2B Intelligence contracts | 1 | 4 | 9 |
| Ronda+ revenue | RM 0.20M | RM 1.35M | RM 3.02M |
| B2B Intelligence revenue | RM 0.35M | RM 4.20M | RM 12.60M |
| Ad-targeting & partnership revenue | RM 0.05M | RM 0.40M | RM 1.10M |
| **Total revenue** | **RM 0.60M** | **RM 5.95M** | **RM 16.72M** |
| **Net position** | **RM (1.80M)** | **RM (0.30M)** | **RM +5.40M** |

***Note on the SOM figure.*** *The RM 72M ARR figure cited in Section 3.3 represents the Year-3 run-rate ceiling of the addressable Klang Valley + two-city market. The financial plan above models a deliberately conservative capture trajectory consistent with B2B sales cycles in Malaysian insurance and energy sectors (typical 6–9 month pilot → 12–18 month enterprise contract).*

## 7.4 Key Milestones

- Month 3 — 8,000 free MAU in Klang Valley. Validation that the product is being shared organically, not just downloaded.

- Month 6 — First B2B Intelligence pilot signed. Validates the thesis that institutional buyers will pay for aggregated rider behavioural data.
- Month 12 — RM 100K monthly recurring revenue between Ronda+ and B2B contracts. Burn-positive on a unit-economics basis even if growth investment continues.
- Month 18 — Operational break-even. Multi-city presence proven. Ready for Series A or organic growth path.
- Month 30 — Cashflow positive. Regional expansion (Indonesia first) feasibility decision.

## 7.5 Capital Required

Ronda is seeking RM 2.4M in seed capital across the 18-month plan, allocated approximately as follows: 55% engineering and product (4 hires plus existing team retention), 22% rider acquisition (referral incentives, content, community partnerships), 13% B2B sales and partnerships (one senior hire plus enablement), 7% data infrastructure (warehouse, anonymisation pipeline, BI tooling), 3% operational overhead. This round is not being raised at UMHackathon; the figure is presented to demonstrate a credible operating plan.

08

# Why This Wins

Three closing arguments.

### The product is shipped

This is not a deck. Ronda is live, integrated end-to-end with ILMU-GLM-5.1, validated against a Pydantic schema with a 100% first-try pass rate, and protected by a three-tier fallback chain ensuring availability under upstream degradation. The /compare endpoint exists in production. The team has already won a national hackathon preliminary round on this build. The judges of the final round will see the working product, not a slide of one.

### The thesis is fundable

Ronda is not a feature. It is a category-creating product in a market with 1.2 million regulated users, no incumbent, three structural moats that compound with every signup, and a B2B revenue layer that maps directly to the Strava and Waze playbooks — both of which produced multi-billion-dollar outcomes from the same pattern of free-app-feeding-paid-data-product. Unit economics work at RM 9 per rider per month. B2B contracts are visible from Year One. The path to RM 16.7 million revenue in three years is conservative.

### The timing is now

The Gig Workers Act 2025 is active law as of 31 March 2026, and 98% of eligible workers are not yet onboarded. Sovereign Malaysian AI infrastructure became production-ready in August 2025. Fuel subsidy rationalisation has put petrol-cost variance into every rider's daily decision. Three regulatory and market forces opened simultaneously. A team with the right combination of product, design, AI integration, and Malaysian context can win this category in the next eighteen months. After that, the data flywheel makes the position structurally hard to displace.

> *1.2 million Malaysians make the same RM-significant decision every morning with no help. We are the help. The market is RM 1.44 billion. The product works today. The only question is whether we move fast enough.*

*Ronda · An AI shift coach for Malaysian gig riders.*

Powered by ILMU-GLM-5.1 · Built on Z.AI and YTL AI Labs sovereign infrastructure.

UMHackathon 2026 · Final Round · Domain 2 · 3 May 2026

A1

# Appendix: Q&A Defense Notes

*For team use during pitch Q&A. Anticipated judge questions and the sharpest answers.*

### "Why won't Grab just build this?"

Because it would harm their marketplace. A platform-built shift coach that honestly says "don't work today, demand is too low" reduces platform supply on a low-revenue day — directly conflicting with the platform's interest in maximum supply availability. The Gig Workers Act 2025 explicitly prohibits unilateral rate manipulation, but it cannot mandate honesty in advisory tools. Our neutrality is structural, not stylistic.

### "What's your CAC, and how is RM 12 realistic?"

Klang Valley rider acquisition runs primarily through community-led seeding, not paid channels. The cost is the RM 10 e-wallet referral incentive plus the marginal cost of community management. We benchmarked this against similar SEA gig-worker acquisition campaigns (food delivery referral programs at launch) which run RM 8–15 per active user. RM 12 is a target not a guarantee; if we miss, we adjust the referral incentive.

### "How is Ronda different from a generic ChatGPT prompt?"

Three things. First, ChatGPT has no live data integration — it cannot see today's weather, today's incentive bulletin, or your remaining BUDI95 quota. Second, ChatGPT confidently hallucinates Malaysian geography in our testing — recommending zones that don't exist or misreading event locations. ILMU-GLM-5.1 is Malay-MMLU-leading and Malaysia-built. Third, Ronda enforces output structure via Pydantic schema — every recommendation has a verdict, windows, projected earnings, and caveats. ChatGPT produces prose that the rider has to re-parse. Ronda produces decisions.

### "Why would an insurance company pay for your data?"

Gig-worker insurance products (motor, accident, income protection) are priced on actuarial assumptions about behaviour the underwriter currently cannot observe. How many hours does a typical rider work? In what conditions do accidents cluster? What's the relationship between platform incentives and risky riding? Today, underwriters guess, and they price conservatively to cover their guesswork. Ronda Intelligence sells them the answer. The first contract is most plausibly with an underwriter launching a Gig Workers Act 2025-compliance product.

### "What if ILMU-GLM-5.1 raises prices or shuts down?"

The Anthropic-compatible API surface means migration to a comparable model is technically straightforward — primarily a configuration change, not a rewrite. Our three-tier fallback chain already proves Ronda functions in degraded mode using only deterministic rules. We would lose Malaysian-context fluency in a forced migration, which is real, but the product would not stop working. We are also

among the first cohort of the ILMU AI Accelerator Programme, which provides USD 1 million in API credits — pricing risk is low through Year 1.

### "This is just a hackathon project. Why should I take the financials seriously?"

The product is shipped. The team has won the preliminary round of a national hackathon judged on production maturity. The financial plan is benchmarked against published comparables (Strava Metro pricing, Waze for Cities contracts, Malaysian SaaS rider-economy revenue per user). Every non-revenue figure in this proposal is sourced from DOSM, PERKESO, MOHR, or named industry publications. We are not asking the judging panel to fund us. We are asking them to evaluate whether the thesis would survive a real diligence conversation. We believe it would.

### "What's the moat in 24 months when someone copies you?"

The proprietary behavioural dataset. By month 24, Ronda will hold the only longitudinal dataset of Malaysian gig-worker decisions outside the platforms. A copycat starting from zero on day one of month 24 would need 24 months to catch up — at which point we are 48 months in. The data flywheel compounds, the trust capital compounds, and the B2B contract revenue funds defensive product investment. The window for entry is closing as this proposal is being read.

*End of document.*